



# High Availability Explained

Maciej Lasyk


11. Sesja Linuksowa

Wrocław, 2014-04-05

“Anything that can go wrong, will go wrong”


Murphy's law

“Anything that can go wrong, will go wrong”



Murphy's law


“Anything that can go wrong, will go wrong”



Murphy's law

An electrical explosion and fire Saturday at a Houston data center operated by The Planet has taken the entire facility offline. The company claimed power to the facility was interrupted when a transformer exploded. Official reports that three walls were blown down causing a fire.

“Anything that can go wrong, will go wrong”



Murphy's law

An electrical explosion and fire Saturday at a Houston data center operated by The Planet has taken the entire facility offline. The company claimed power to the facility was interrupted when a transformer exploded. Official reports that three walls were blown down causing a fire.

Three walls of the electrical equipment room on the first floor blew several feet from their original position, and the underground cabling that powers the first floor of H1 was destroyed.

# High Availability is in the eye of the beholder

# High Availability is in the eye of the beholder

CEO: we don't loose sales

# High Availability is in the eye of the beholder

CEO: we don't loose sales

Sales: we can extend our offer basing on HA level



# High Availability is in the eye of the beholder

CEO: we don't loose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

# High Availability is in the eye of the beholder

CEO: we don't loose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

Developers: we can be proud – our services are working ;)

# High Availability is in the eye of the beholder

CEO: we don't lose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

Developers: we can be proud – our services are working ;)

System engineers: we can sleep well (and f\*ck, we love to!)

# High Availability is in the eye of the beholder

CEO: we don't loose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

Developers: we can be proud – our services are working ;)


System engineers: we can sleep well (and fsck, we love to!)

Technical support: no calls? Back to WoW then.. ;)

# So how many 9's?

# So how many 9's?

SIMPLY EXPLAINED



# So how many 9's?

Monthly: 1 hour of outage means  $100\% - 0.13888 \approx \mathbf{99.86112}$  of availability

# So how many 9's?

**Monthly:** 1 hour of outage means  $100\% - 0.13888 \approx \mathbf{99.86112}$  of availability

**Yearly:** 1 hour of outage means  $100\% - 0.01142 \approx \mathbf{99.98858}$  of availability



# So how many 9's?

**Monthly:** 1 hour of outage means  $100\% - 0.13888 \approx \mathbf{99.86112}$  of availability

**Yearly:** 1 hour of outage means  $100\% - 0.01142 \approx \mathbf{99.98858}$  of availability

<b>Availability</b>	<b>Downtime (year)</b>	<b>Downtime (month)</b>
90% ("one nine")	36.5 days	72 hours
95%	18.25 days	36 hours
97%	10.96 days	21.6 hours
98%	7.30 days	14.4 hours
99% ("two nines")	3.65 days	7.2 hours
99.5%	1.83 days	3.6 hours
99.8%	17.52 hours	86.23 minutes
99.9% ("three nines")	4.38 hours	21.56 minutes
99.99 ("four nines")	52.56 minutes	4.32 minutes
<b>99.999 ("five nines")</b>	<b>5.26 minutes</b>	<b>25.9 seconds</b>

# So how many 9's?




Figure 1-3 Cost of availability as opposed to loss due to availability

<https://jazz.net/wiki/bin/view/Deployment/HighAvailability>

# HA terminology

**RPO:** Recovery Point Objective; how much data can we loose?

# HA terminology

**RPO:** Recovery Point Objective; how much data can we loose?

**RTO:** Recovery Time Objective; how long does it take to recover?


# HA terminology

**RPO:** Recovery Point Objective; how much data can we loose?

**RTO:** Recovery Time Objective; how long does it take to recover?

**MTBF:** Mean-Times-Between-Failures; time between failures

(density fnc -> reliability fnc)  $\int_0^{\infty} f(t) dt = 1.$



Time Between Failures = { down time - up time }

[https://en.wikipedia.org/wiki/Mean\\_time\\_between\\_failures](https://en.wikipedia.org/wiki/Mean_time_between_failures)

# HA terminology

**SLA:** Service Level Agreement;  
formal definitions (customer <-> provider)

# HA terminology

**SLA:** Service Level Agreement;  
formal definitions (customer <-> provider)

**OLA:** Operational Level Agreement; definitions within organization;  
help us keeping provided SLAs

# SLAs..

So what is written in **SLAs**?

<b>Availability</b>	<b>Downtime (year)</b>	<b>Downtime (month)</b>
90%	36.5 days	72 hours
95%	18.25 days	36 hours
97%	10.96 days	21.6 hours
98%	7.30 days	14.4 hours
99%	3.65 days	7.2 hours
99.5%	1.83 days	3.6 hours
99.8%	17.52 hours	86.23 minutes
99.9%	4.38 hours	21.56 minutes
99.99	52.56 minutes	4.32 minutes
<b>99.999</b>	<b>5.26 minutes</b>	<b>25.9 seconds</b>



# SLAs..

So what is written in **SLAs**?

<b>Availability</b>	<b>Downtime (year)</b>	<b>Downtime (month)</b>
90%	36.5 days	72 hours
95%	18.25 days	36 hours
97%	10.96 days	21.6 hours
98%	7.30 days	14.4 hours
99%	3.65 days	7.2 hours
<b>99.5% (EC2, EBS)</b>	1.83 days	3.6 hours
99.8%	17.52 hours	86.23 minutes
<b>99.9% (SoftLayer, IBM)</b>	4.38 hours	21.56 minutes
<b>99.99 (OVH ded. cloud)</b>	52.56 minutes	4.32 minutes
<b>99.999</b>	<b>5.26 minutes</b>	<b>25.9 seconds</b>

<http://aws.amazon.com/ec2/sla/>

<http://www.softlayer.com/about/service-level-agreement>

<https://www.ovh.com/us/dedicated-cloud/security-and-sla.xml>

# SLAs..

Availability mentioned in **SLAs** are only **goals** of service provider

Usually when it's **not met** than company **pays off** the fees

# SLAs..

## **Hetzner?**

“We guarantee an annual average of 99% network availability”

“For indirect damages and loss of profits, we are liable only in cases of intentional or gross negligence. In this case we are liable only for the contract-typical predictable damage, a maximum of 100% of the annually fee.”

<http://www.hetzner.de/en/hosting/legal/agb>

# SLAs..

## SLA PACKAGES

### Leaseweb?

(yup - megaupload)

- no %s
- best effort
- \$\$\$ for faster response times

	Basic	Bronze	Silver	Gold	Platinum
Free phone support	24/7	24/7	24/7	24/7	24/7
Response time	24 hrs	4 hrs	2 hrs	1 hr	30 min
Hardware replacement**	24 hrs	4 hrs	3 hrs	2 hrs	2 hrs
Self Service Center	V	V	V	V	V
Advanced support hourly rate	€99 / \$129	€89 / \$119	€79 / \$99	€69 / \$89	€59 / \$79
Advanced support included*	-	30 min	60 min	90 min	120 min
Price/month	Included with all dedicated servers	€29 / \$39	€49 / \$59	€79 / \$99	€119 / \$129

<http://www.leaseweb.com/en/support/all-about/sla>

# How deep is this hole?

app layer (core, db, cache)

data storage

operating system

hardware

networking

location

So we would like to achieve 99,9999% which is about 30s of downtime per year

# How deep is this hole?

app layer (core, db, cache)

data storage

operating system

hardware

networking

location

Even Proof of Concept is very hard to provide: 2.5s of downtime per layer monthly!

# How deep is this hole?

app layer (core, db, cache)

data storage

operating system

hardware

networking

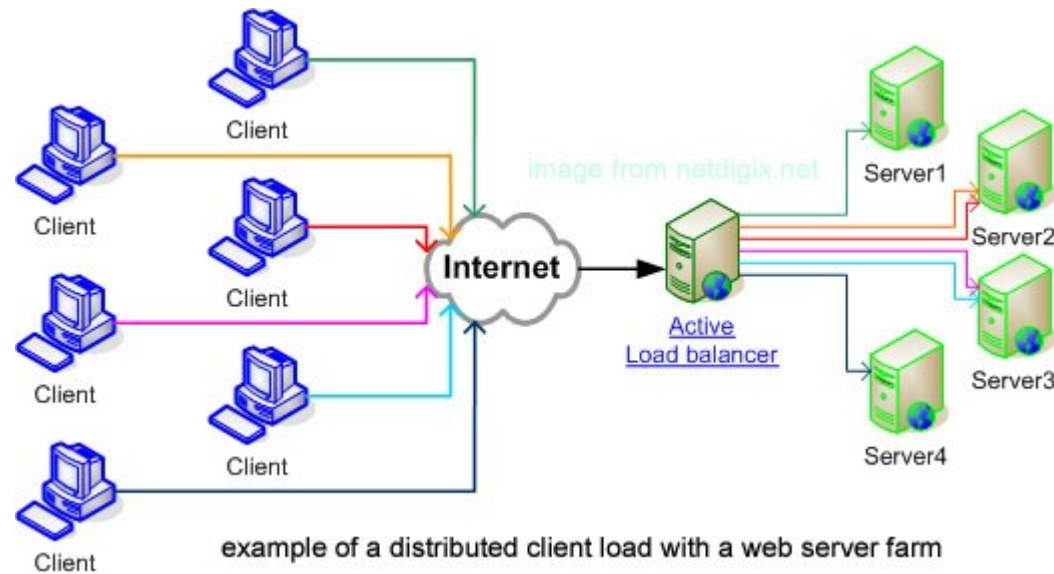
location

AOL has 99,999%!

<http://highscalability.com/blog/2014/2/17/how-the-aolcom-architecture-evolved-to-99999-availability-8.html>

# Load-balancing and failover

LB:




<http://www.netdigix.com/linux-loadbalancing.php>



# Load-balancing and failover

Failover:



<http://www.simplefailover.com/>

# LB – 4<sup>th</sup> layer or 7<sup>th</sup>?

## 4<sup>th</sup> layer:

- high performance
- just do the LB work!
- reliable
- scalable


## 7<sup>th</sup> layer:

- low cost
- good for quickfixes / patches
- not that scalable
- low performance
- complex codebase
- custom code for protocols
- cookies? what about memcache..

# Disaster Recovery




# Disaster Recovery



<http://disasterrecovery.starwindsoftware.com/planning-disaster-recovery-for-virtualized-environments>

# Disaster Recovery




<http://disasterrecovery.starwindsoftware.com/planning-disaster-recovery-for-virtualized-environments>

**Hot site:** active synchronization, could be serving services. Cost can be high

**Warm site:** periodical synchronization, DR tests needed. Low costs

**Cold site:** Nothing here – just echo and some place to spin services; nightmare

# Disaster Recovery




<http://disasterrecovery.starwindsoftware.com/planning-disaster-recovery-for-virtualized-environments>

**Hot site:** active synchronization, could be serving services. Cost can be high

**So maybe use “hot” as production and global LB? win-win scenario**

# Planning for failure

*ENTEPRISE ARCHITECTURE MADE EASY*



*PART 1: DON'T MESS WITH THE GORY DETAILS*

# Planning for failure

## Everything starts here - DNS:

- keep TTLs low (300s). Can't make under 60min? That's bad!
- check SLA of DNS servers (dnsmadeeasy.com history)
- what do you know about DNSes?
  - zero downtime here is a must!
  - this can be achieved with complicated network abracadabra
  - remember what 99.9999% means?
- round robin is a load – balancer but **without failover!**
- GSLB – killed by OS/browser/srvs cache'ing  
(GlobalServerLoadBalancing)
- GlobalIP (SoftLayer etc) – workaround for GSLB via routing



# Planning for failure

## E-mail servers:

- it's simple as MX records (**delivering**)
- it's almost simple as complicated system of SMTP servers (**sending**)
- it's not that simple when IMAP locking over DFS (**reading**)

```
5 gmail-smtp-in.l.google.com.  
10 alt1.gmail-smtp-in.l.google.com.  
20 alt2.gmail-smtp-in.l.google.com.  
30 alt3.gmail-smtp-in.l.google.com.  
40 alt4.gmail-smtp-in.l.google.com.
```

When MXing – **watch the spam!**

# Planning for failure

## WEB servers:

- it's simple as some frontend loadbalancer
- did you **really** stick user session to particular server? Memcache!
- LB balancing algorithm
- how many LBs?
- what if LB goes down?

# Planning for failure

## DB servers:

- it's.. not that simple
- replication (master – master? App should be aware..)
- replication ring? Complicated, works, but in case of failure...
- let's talk about MySQL:
  - NoSPOF solution: MySQL cluster
  - MySQL Galera cluster – synch, active-active multi-master
  - master – master – simply works
  - MySQL fabric – HA + sharding; use with large farms
  - Failover? Matsunobu Yoshinori `mysql-master-ha`
  - MySQL utilities (<http://www.clusterdb.com/mysql/mysql-utilities-webinar-qa-replay-now-available/>)

# Planning for failure

## DB servers:

Matsunobu Yoshinori mysql-master-ha

<https://code.google.com/p/mysql-master-ha/>

“I have heard a couple of cases where AWS users use MHA instead of RDS because RDS takes much longer downtime on slave promotion (more than 5 minutes usually, because standby database is not running).

I'm surprised AWS users care about a few minutes of downtime..”

# Planning for failure

## Caching servers:

- this is cache for God's sake – why would we use HA here?
- just use proper architecture like... redundancy.

# Planning for failure

## Caching servers:

- this is cache for God's sake – why would we use HA here?
- just use proper architecture like... redundancy.

## Load – balancers:

- remember about failovering IP addresses!

# Planning for failure

## Caching servers:

- this is cache for God's sake – why would we use HA here?
- just use proper architecture like... redundancy.

## Load – balancers:

- remember about failovering IP addresses!

## Storage – DFSes:

- GlusterFS – we'll see it in action in a minute
- NFS? Could be – over some SAN / NAS (high cost solution)
- CephFS – just like GlusterFS – it's great and does the work
- DRBD – lower level, does the work on block – device layer – slow...

# Planning for failure

## GlusterFS:

- low cost (could be..)
- distributed volumes
- replicated volumes
- striped volumes
- and...
  - distributed – striped volumes
  - distributed – replicated volumes
  - distributed – striped – replicated volumes
- sound good? :)



# Planning for failure

## GlusterFS: replicated volumes vs Geo-replication

- replicated:

- mirrors data
- provides HA
- synch – replication

- Geo-replication:

- mirrors data across geo – distributed clusters
- ensures backing up data for DR
- asynch – replica (periodic checks)

# Planning for failure

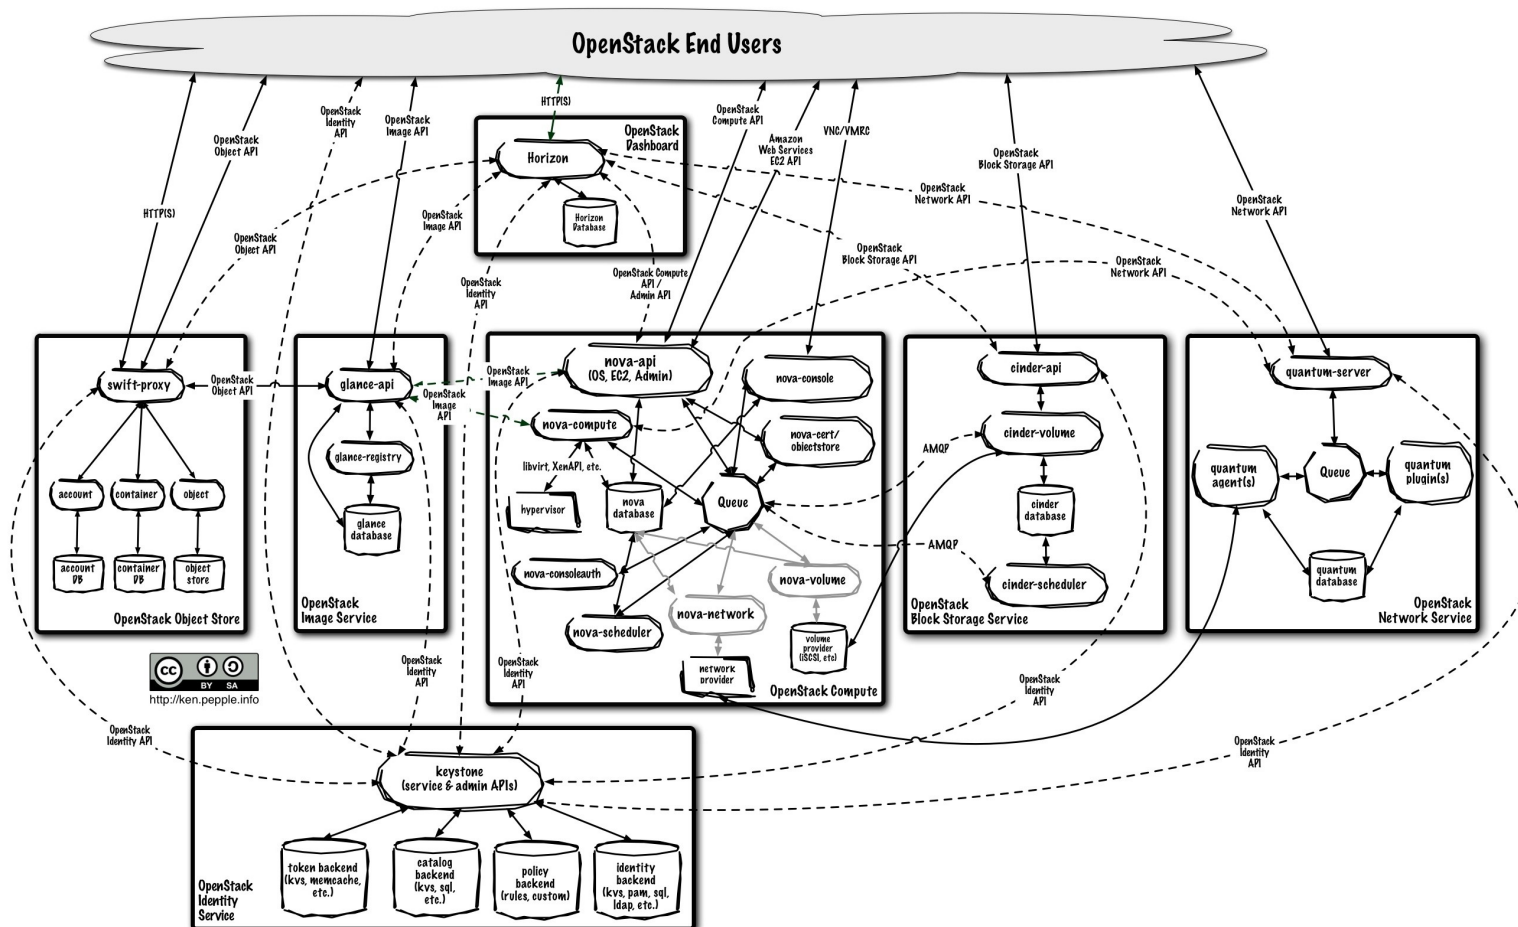
HA for virtualization solutions?

- it's really complicated, like...

# Planning for failure

## HA for virtualization solutions?


- it's really complicated, like...



# Planning for failure

HA for virtualization solutions?


- but it could be done simpler...



# Planning for failure

## HA for virtualization solutions?

- it could be done simpler with containers!
- containers are very light
- deploy time from bare is tiny
- management is very easy
- resources throttling via cgroups





docker

# Tools

The most important tool would be the conclusion from the picture below:

# Tools



The most important tool would be the conclusion from the picture below:



LET THE CLOUDS MAKE YOUR LIFE EASIER

# Tools

The most important tool would be the conclusion from the picture below:



LET THE CLOUDS MAKE YOUR LIFE EASIER





# Tools

- DNS: roundrobin, GSLB, low ttls, globalIP

# Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx

# Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx
- Failover (statefull services):
  - IP: KeepAlived + sysctl

# Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx
- Failover (statefull services):
  - IP: KeepAlived + sysctl
  - Managing: pacemaker (manager) + corosync (message'ing)

# Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx
- Failover (statefull services):
  - IP: KeepAlived + sysctl
  - Managing: pacemaker (manager) + corosync (message'ing)
- (almost) All-In-One: Linux Virtual Server

# Turn on HA thinking!

Main goal of HA? **Improve user experience!**

- keep the app **fully functional**
- keep the app **resistant and tolerant** to faults
- provide method for a **successful audit**
- **sleep well** (anyone awake?) ;)

# Thank you :)

## High Availability Explained

Maciej Lasyk

11. Sesja Linuksowa

2014-04-05, Wrocław

<http://maciek.lasyk.info/sysop>

[maciek@lasyk.info](mailto:maciek@lasyk.info)

@docent-net