


Orchestrating Docker containers at scale

Maciej Lasyk
12 Sesja Linuksowa
Wrocław 2015-04-18



Join Fedora Infrastructure!

- learn Ansible
- learn Docker with Fedora Dockerfiles

<http://fedoraproject.org/en/join-fedora>

Quick survey

How many of you...

Knows what Docker is?

Played with Docker?

Runs it on production?

Quick survey

How many of you...

Knows what Docker is?

Played with Docker?

Runs it on production?

Quick survey

How many of you...

Knows what Docker is?

Played with Docker?

Runs it on production?

Quick survey

How many of you...

Knows what Docker is?

Played with Docker?

Runs it on production?

Why use Docker?

With Docker we can solve many problems

- “it works on my machine”
- reducing build & deploy time
- Infrastructure configuration spaghetti – automation!
- Libs dependency hell
- Cost control and granularity

Why use Docker?

With Docker we can solve many problems

- “it works on my machine”
- reducing build & deploy time
- Infrastructure configuration spaghetti – automation!
- Libs dependency hell
- Cost control and granularity

Why use Docker?

With Docker we can solve many problems

- “it works on my machine”
- reducing build & deploy time
- Infrastructure configuration spaghetti – automation!
- Libs dependency hell
- Cost control and granularity

Why use Docker?

With Docker we can solve many problems

- “it works on my machine”
- reducing build & deploy time
- Infrastructure configuration spaghetti – automation!
- Libs dependency hell
- Cost control and granularity

Why use Docker?

With Docker we can solve many problems

- “it works on my machine”
- reducing build & deploy time
- Infrastructure configuration spaghetti – automation!
- **Libs dependency hell**
- Cost control and granularity

Why use Docker?

With Docker we can solve many problems

- “it works on my machine”
- reducing build & deploy time
- Infrastructure configuration spaghetti – automation!
- Libs dependency hell
- Cost control and granularity

Docker – what is it?



“automates the deployment of any application as a lightweight, portable, self-sufficient container that will run virtually anywhere”

Java's promise: Write Once. Run Anywhere.

Java's promise: Write Once. Run Anywhere.



Even on Windows now!

<https://blog.docker.com/2014/10/docker-microsoft-partner-distributed-applications/>

Is Docker is lightweight?

Is Docker is lightweight?


=====				
Package	Arch	Version	Repository	Size
=====				
Installing:				
docker-io	x86_64	1.3.0-1.fc20	updates	4.3 M

Is Docker is lightweight?


Package	Arch	Version	Repository	Size
Installing: docker-io	x86_64	1.3.0-1.fc20	updates	4.3 M

Package	Arch	Version	Repository	Size
Installing: docker-io	x86_64	1.5.0-2.fc21	updates	26 M

VMs vs. Containers



Docker – how it works?



Docker – how it works?

→ LXC & libcontainer

→ control groups

→ kernel namespaces

→ layered filesystem

→ no more AUFS (perf sucks)

→ devmapper thin provisioning & loopback mounts

→ OverlayFS!

Docker – how it works?

- LXC & libcontainer
- **control groups**
- kernel namespaces
- layered filesystem
 - no more AUFS (perf sucks)
 - devmapper thin provisioning & loopback mounts
 - OverlayFS!

Docker – how it works?

- LXC & libcontainer
- control groups
- **kernel namespaces**
- layered filesystem
 - no more AUFS (perf sucks)
 - devmapper thin provisioning & loopback mounts
 - OverlayFS!

Docker – how it works?

- LXC & libcontainer
- control groups
- kernel namespaces
- **layered filesystem**
 - no more AUFS (perf sucks)
 - devmapper thin provisioning & loopback mounts
 - OverlayFS!

Docker – how it works?

- LXC & libcontainer
- control groups
- kernel namespaces
- **layered filesystem**
 - **no more AUFS (perf sucks)**
 - devmapper thin provisioning & loopback mounts
 - OverlayFS!

Docker – how it works?

- LXC & libcontainer
- control groups
- kernel namespaces
- **layered filesystem**
 - no more AUFS (perf sucks)
 - **devmapper thin provisioning & loopback mounts**
 - OverlayFS!

Docker – how it works?

- LXC & libcontainer
- control groups
- kernel namespaces
- **layered filesystem**
 - no more AUFS (perf sucks)
 - devmapper thin provisioning & loopback mounts
 - **OverlayFS!**

control groups (cgroups)

Control Groups provide a mechanism for aggregating/partitioning sets of tasks, and all their future children, into hierarchical groups with specialized behavior

control groups (cgroups)

- grouping processes
- allocating resources to particular groups
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

- grouping processes
- **allocating resources to particular groups**
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

- grouping processes
- allocating resources to particular groups
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

- grouping processes
- allocating resources to particular groups
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

- grouping processes
- allocating resources to particular groups
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

- grouping processes
- allocating resources to particular groups
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

- grouping processes
- allocating resources to particular groups
 - memory
 - network
 - CPU
 - storage bandwidth (I/O throttling)
 - device whitelisting

control groups (cgroups)

little demo?

Kernel Namespaces

Providing a unique views of the system for processes.

- PID – PIDs isolation
- NET – network isolation (via virt-ifaces; demo)
- IPC – won't use this
- MNT – chroot like; deals w/mountpoints
- UTS – deals w/hostname

Kernel Namespaces

Providing a unique views of the system for processes.

- PID – PIDs isolation
- NET – network isolation (via virt-ifaces; demo)
- IPC – won't use this
- MNT – chroot like; deals w/mountpoints
- UTS – deals w/hostname

Kernel Namespaces

Providing a unique views of the system for processes.

- PID – PIDs isolation
- NET – network isolation (via virt-ifaces; demo)
- IPC – won't use this
- MNT – chroot like; deals w/mountpoints
- UTS – deals w/hostname

Kernel Namespaces

Providing a unique views of the system for processes.

- PID – PIDs isolation
- NET – network isolation (via virt-ifaces; demo)
- IPC – won't use this
- MNT – chroot like; deals w/mountpoints
- UTS – deals w/hostname

Kernel Namespaces

Providing a unique views of the system for processes.

- PID – PIDs isolation
- NET – network isolation (via virt-ifaces; demo)
- IPC – won't use this
- MNT – chroot like; deals w/mountpoints
- UTS – deals w/hostname

Kernel Namespaces

Providing a unique views of the system for processes.

- PID – PIDs isolation
- NET – network isolation (via virt-ifaces; demo)
- IPC – won't use this
- MNT – chroot like; deals w/mountpoints
- UTS – deals w/hostname

Kernel Namespaces

little demo?

OverlayFS

- hell fast (you'll see)
- page cache sharing
- finally in upstream kernel (in rhel from 7.2)
- finally supported by docker (-s overlay)
- SELinux not there yet (but will be)

OverlayFS

- hell fast (you'll see)
- **page cache sharing**
- finally in upstream kernel (in rhel from 7.2)
- finally supported by docker (-s overlay)
- SELinux not there yet (but will be)

OverlayFS

- hell fast (you'll see)
- page cache sharing
- **finally in upstream kernel (in rhel from 7.2)**
- finally supported by docker (-s overlay)
- SELinux not there yet (but will be)

OverlayFS


- hell fast (you'll see)
- page cache sharing
- finally in upstream kernel (in rhel from 7.2)
- **finally supported by docker (-s overlay)**
- SELinux not there yet (but will be)

OverlayFS

- hell fast (you'll see)
- page cache sharing
- finally in upstream kernel (in rhel from 7.2)
- finally supported by docker (-s overlay)
- SELinux not there yet (but will be)

OverlayFS


Container Create/Destroy Times



OverlayFS

Docker Page Cache Usage Test

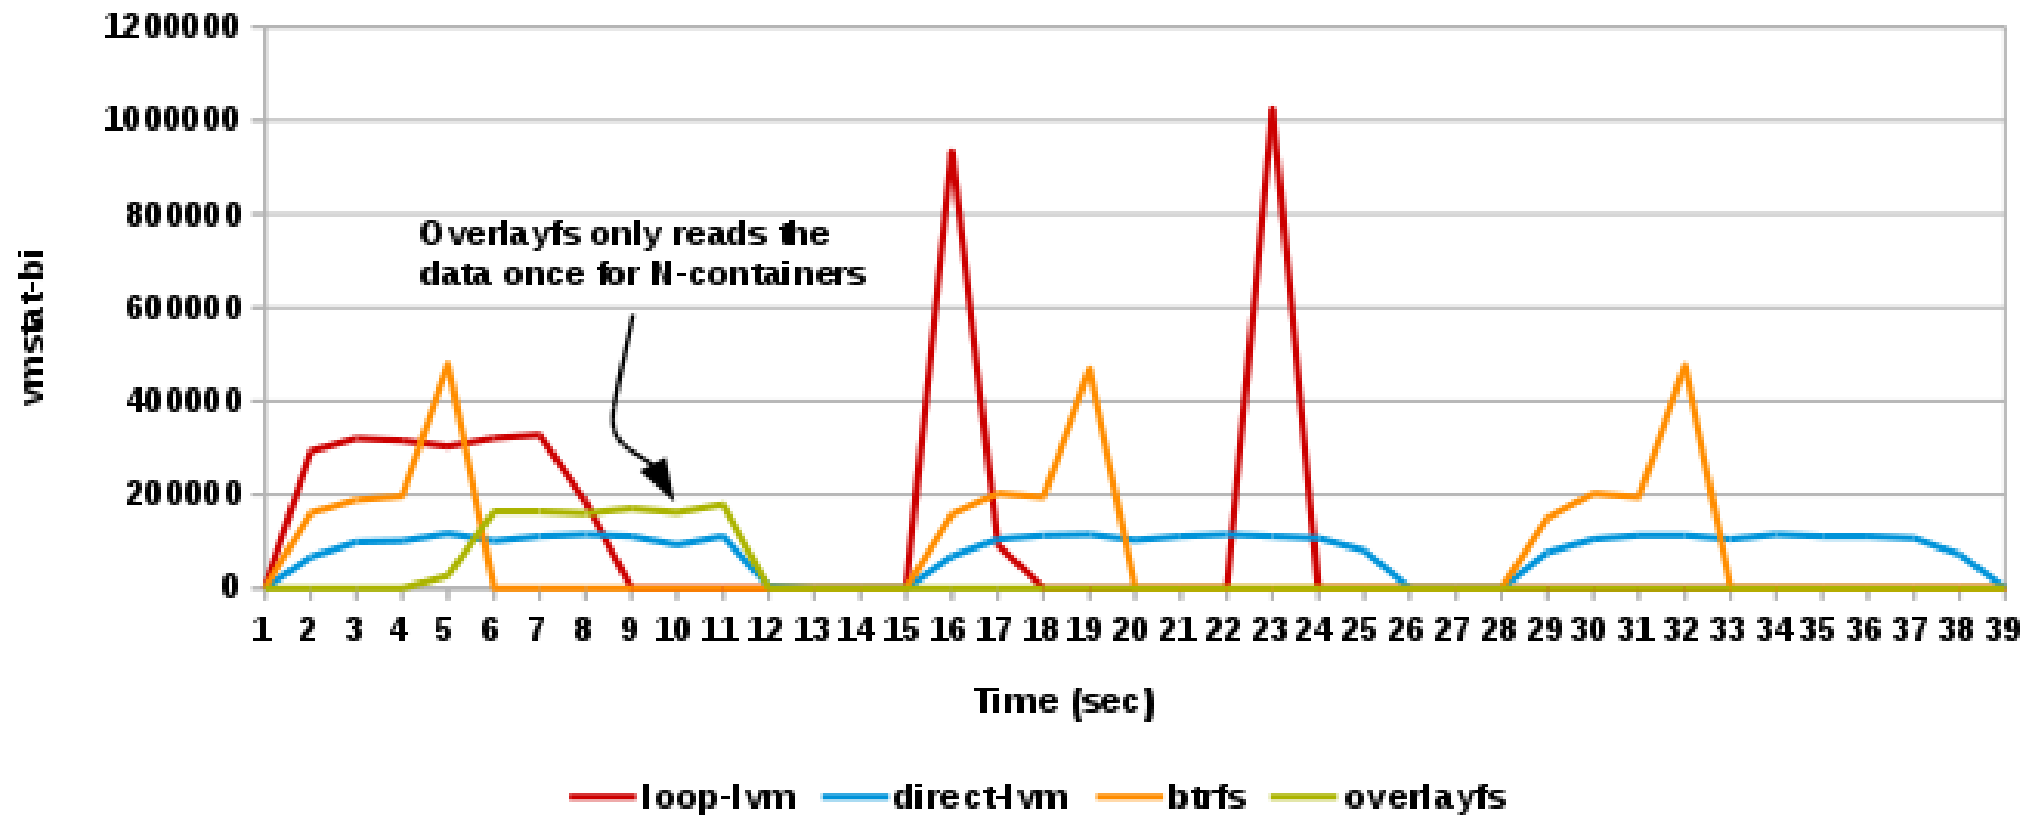
docker-1.1 + 3.17-rc1



OverlayFS

Docker Page Cache Usage Test

docker-1.1 + 3.17-rc1



OverlayFS

little demo?

Linux containers equation

Linux Containers = namespaces + cgroups + storage

Docker – concepts

Images

- read only
- act as templates

Dockerfile

- like a makefile
- commands order & cache'ing
- extends the base image
- results in a new image

Containers: instances running apps

Docker – concepts

Images

- read only
- act as templates

Dockerfile

- like a makefile
- commands order & cache'ing
- extends the base image
- results in a new image

Containers: instances running apps

Docker – concepts

Images

- read only
- act as templates

Dockerfile

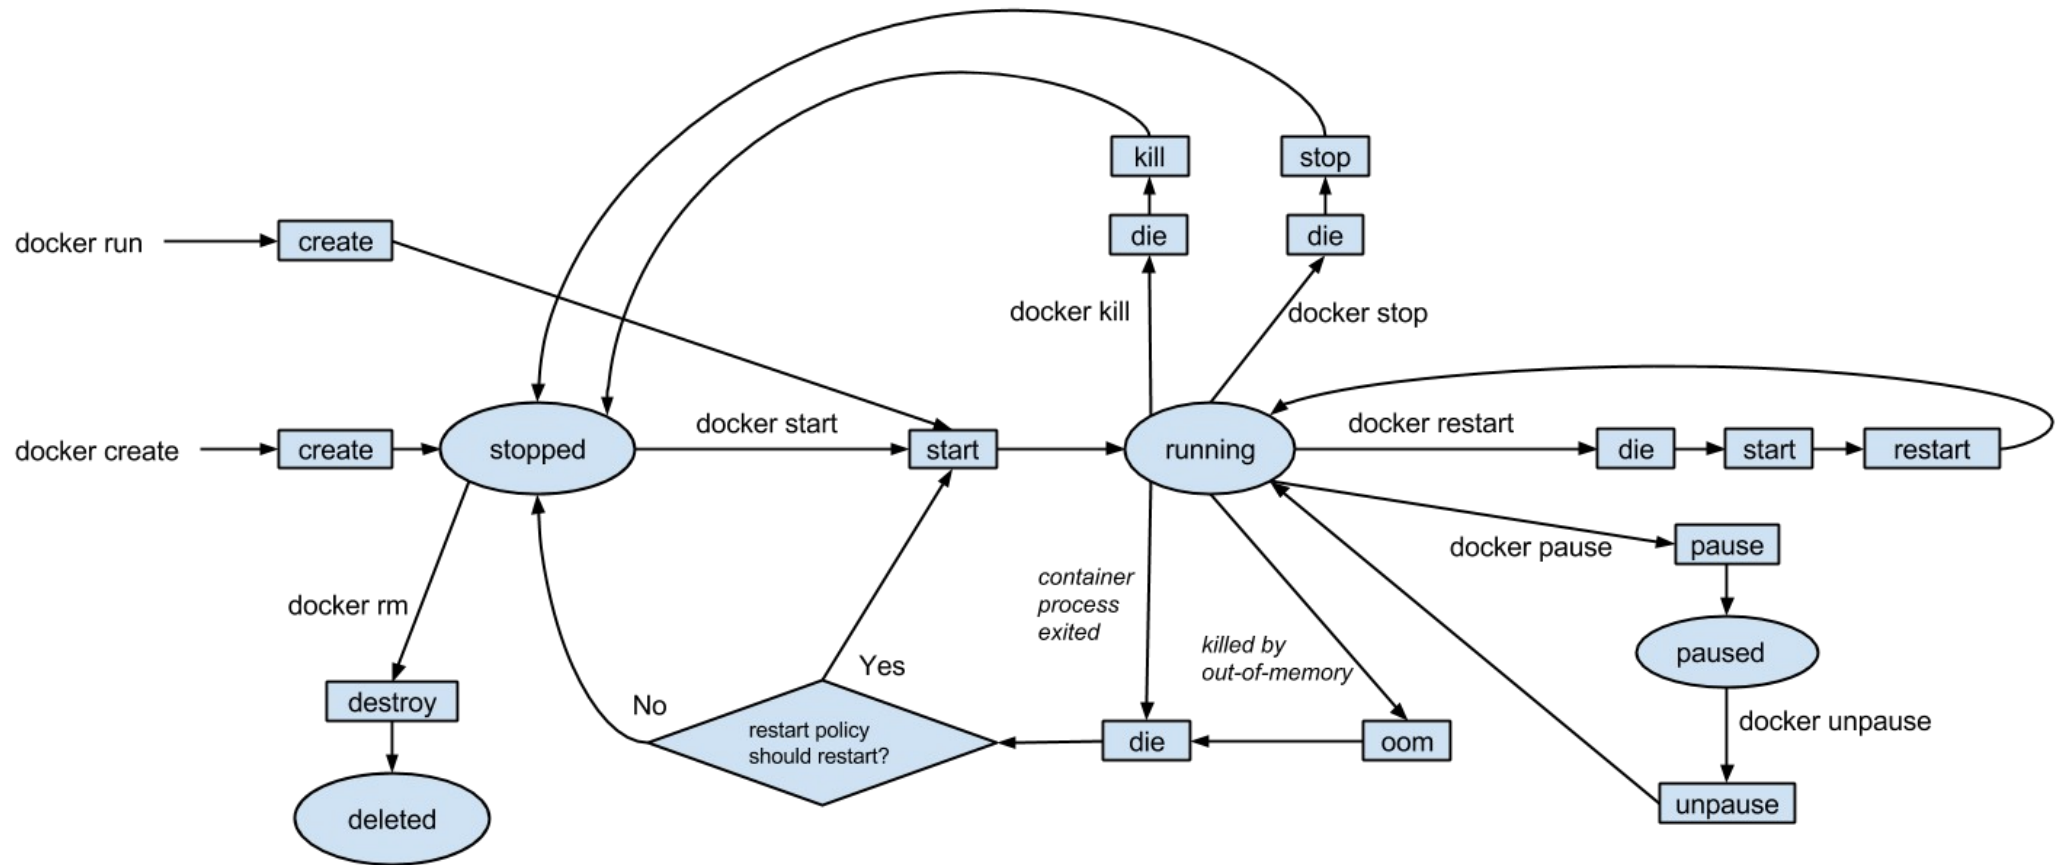
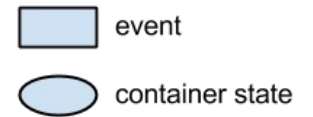
- like a makefile
- commands order & cache'ing
- extends the base image
- results in a new image

Containers: instances running apps

Docker – concepts

`dockerfile + base image = docker container`

Docker - events



<http://gliderlabs.com/blog/2015/04/14/docker-events-explained/>

Dockerfile

FROM fedora

MAINTAINER scollier <scollier@redhat.com>

RUN yum -y update && yum clean all

RUN yum -y install nginx && yum clean all


RUN echo "daemon off;" >> /etc/nginx/nginx.conf

RUN echo "nginx on Fedora" > /srv/www/index.html

EXPOSE 80

CMD ["/usr/sbin/nginx"]

Docker – registry



Docker – registry

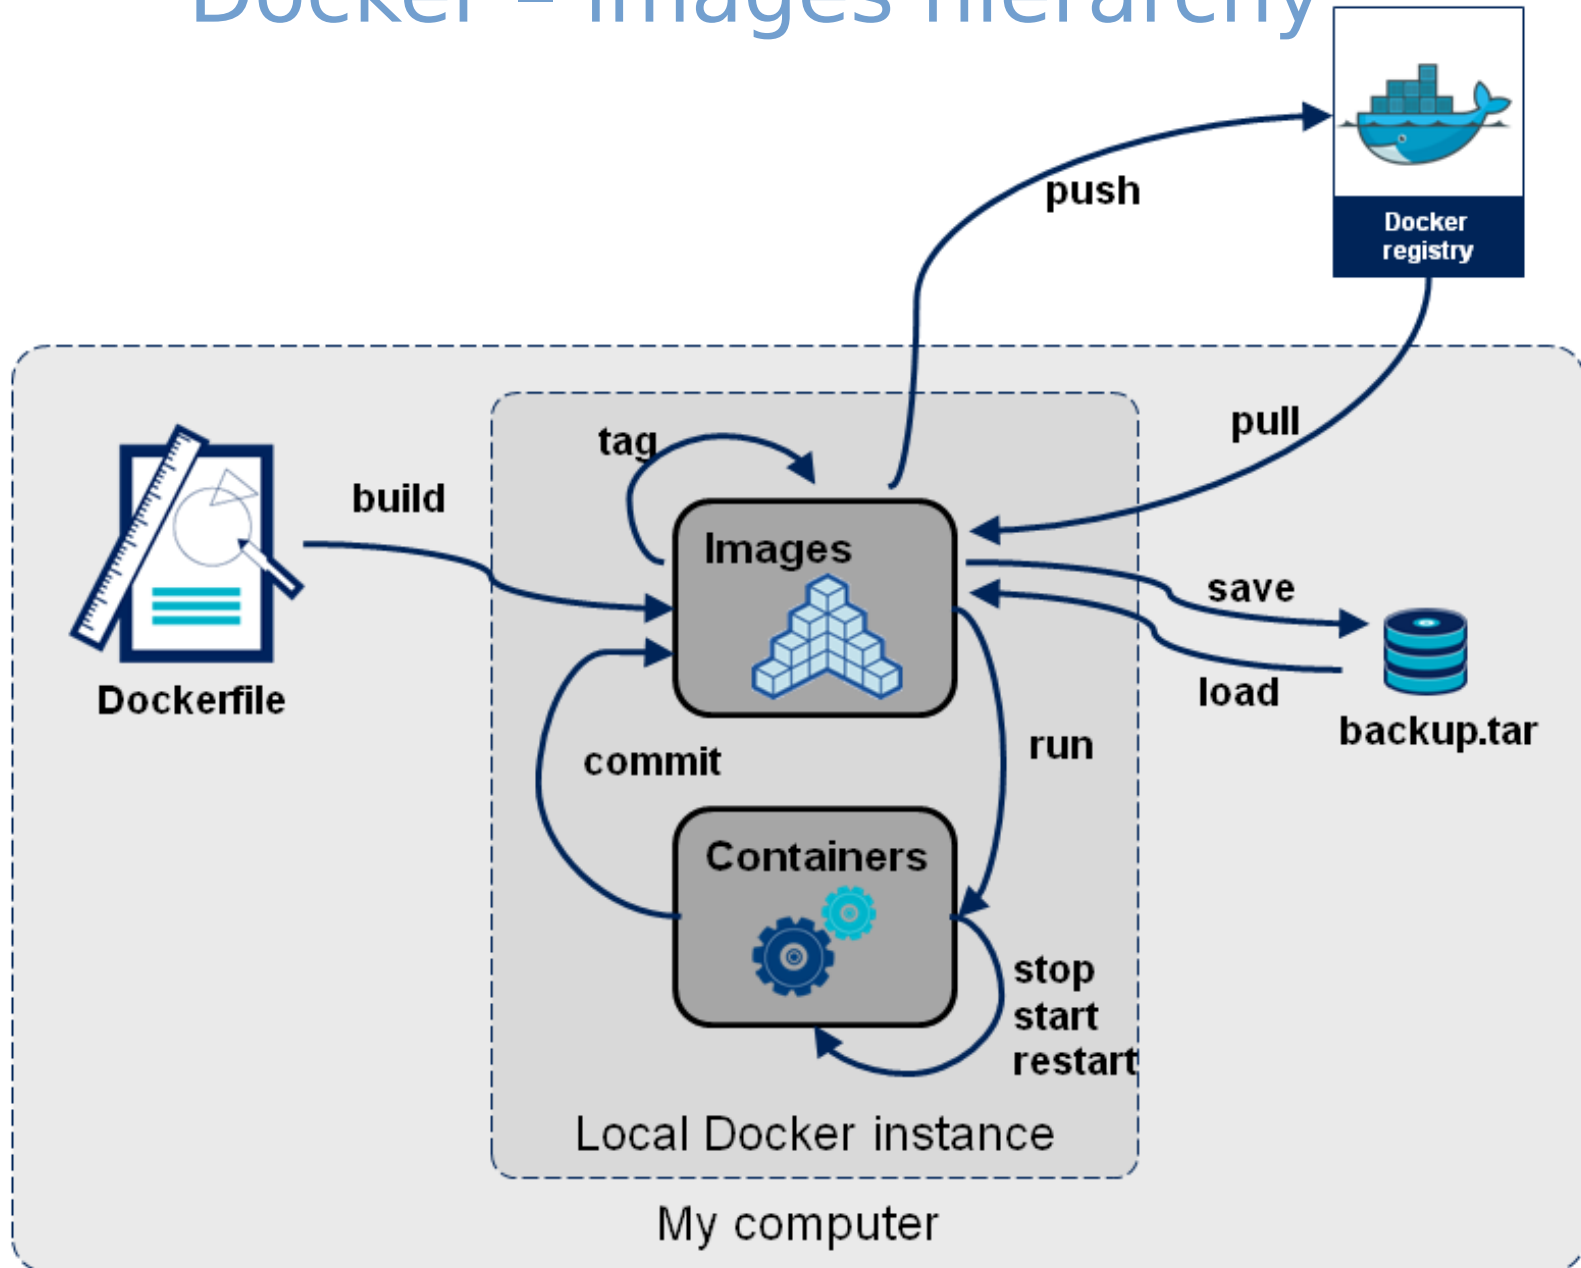
- git like semantics
- pull, push, commit
- private and public registry
- <https://github.com/dotcloud/docker-registry>
- yum install docker-registry

```
$ docker pull
```

```
$ docker push
```

```
$ docker commit
```

Docker – images hierarchy



Docker – images hierarchy

base image

-> child image

-> grandchild image

Docker – images hierarchy

base image

-> child image

-> grandchild image

Git's promise: Tiny footprint with
lightning fast performance

Docker – security

- Isolation via kernel namespaces
- Each container gets own network stack
- Control groups for resources limiting
- Additional layer of security: SELinux / AppArmor / GRSEC

f20 policy: <https://git.fedorahosted.org/cgit/selinux-policy.git/tree/docker.te?h=f20-contrib>

What's there?

```
seinfo -t -x | grep docker
```

```
sesearch -A -s docker_t (and the rest)
```

```
or just unpack docker.pp with semodule_unpackage
```


Docker – security

Docker has changed its security status to
It's complicated

Docker – security



Docker – use cases



CI Stack

Docker – use cases

- Continuous Integration
- local dev
 - with Docker it's easy to standardize envs
- deployment (rolling updates (e.g. w/Ansible))
- testing
 - unit testing of any commit on dedicated env
 - don't worry about cleaning up after testing
 - paralleled tests across any machines

Docker – use cases

- version control system for apps
- microservices
 - Docker embraces granularity
 - Services can be deployed independently and faster
 - paralleled tests across any machines
- continuous delivery
- PaaS

Docker – history

- 2013-01: dotCloud worked on own PaaS (Python based)
- 2013-03: Docker went public (AUFS, LXC)
- middle 2013: Red Hat joined, devmapper, SELinux
- late 2013: removed LXC, rewritten in Go
- 2014-02: stable 1.0

Docker – history

- 2013-01: dotCloud worked on own PaaS (Python based)
- **2013-03: Docker went public (AUFS, LXC)**
- middle 2013: Red Hat joined, devmapper, SELinux
- late 2013: removed LXC, rewritten in Go
- 2014-02: stable 1.0

Docker – history

- 2013-01: dotCloud worked on own PaaS (Python based)
- 2013-03: Docker went public (AUFS, LXC)
- **middle 2013: Red Hat joined, devmapper, SELinux**
- late 2013: removed LXC, rewritten in Go
- 2014-02: stable 1.0

Docker – history

- 2013-01: dotCloud worked on own PaaS (Python based)
- 2013-03: Docker went public (AUFS, LXC)
- middle 2013: Red Hat joined, devmapper, SELinux
- **late 2013: removed LXC, rewritten in Go**
- 2014-02: stable 1.0

Docker – history

- 2013-01: dotCloud worked on own PaaS (Python based)
- 2013-03: Docker went public (AUFS, LXC)
- middle 2013: Red Hat joined, devmapper, SELinux
- late 2013: removed LXC, rewritten in Go
- 2014-02: stable 1.0

Docker – popularity

Compare

Search terms ▾

docker

Search term

lxc

Search term

kvm

Search term

xen


Search term

+ Add term

Interest over time ?

☒ News headlines

☐ Forecast ?



</>

Docker – popularity

Compare

Search terms ▼

docker

Search term

lxc

Search term

kvm

Search term

xen

Search term


vmware

Search term

Interest over time ?

☒ News headlines

☐ Forecast ?



</>

Orchestration at scale w/Docker



Orchestration at scale w/Docker

This might be a little problem

Orchestration at scale w/Docker

Service Providers



Operating Systems




Configuration Management



Big Data



Service Discovery



Dev Tools



Official Repositories



Orchestration



System Integrators



Orchestration at scale w/Docker



	Big Data	Cloud Platform	IaaS	Data Center OS	Docker OS	Docker Mgmt.	PaaS	Orch. Config Mgmt.
Ansible & Docker								x
Amazon EC2 & Docker		x						
Apache Brooklyn & Docker								x
Apache Hadoop & Docker	x							
Apache Storm & Docker	x							
AppScale & Docker							x	
Atomic Hosts & Docker					x			
Chef & Docker								x
Clocker & Docker								x
Cloud Foundry & Docker	x						x	
CloudStack & Docker			x					
CoreOS & Docker					x			
Deis & Docker							x	
Decker & Docker							x	
Docker & Docker			x		x	x	x	x
Dokku & Docker							x	
Eucalyptus & Docker			x					
Flynn & Docker							x	
Google Compute Platform & Docker		x						
IBM Bluemix & Docker	x						x	
Kubernetes & Docker			x			x	x	x
Mesos, Mesosphere & Docker	x			x		x	x	x
Microsoft Azure & Docker		x						
OpenCamp & Docker		x	x			x	x	
OpenShift & Docker							x	
OpenStack & Docker			x					
Panamax & Docker						x		
Puppet & Docker								x
SaltStack & Docker							x	x
Shipyards & Docker						x		
Stackato & Docker							x	
Tsuru & Docker							x	
VMware & Docker			x					

<http://www.cloudssky.com/en/blog/Docker-Is-Not-Enough>

To Paas or not to PaaS?



To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- installing PaaS might be easy
- operating PaaS will be tough
- so is operating your apps tough enough to move?
- private PaaS or not?
- is your app ready for PaaS?
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- **it will rather clone them :)**
- installing PaaS might be easy
- operating PaaS will be tough
- so is operating your apps tough enough to move?
- private PaaS or not?
- is your app ready for PaaS?
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- **installing PaaS might be easy**
- operating PaaS will be tough
- so is operating your apps tough enough to move?
- private PaaS or not?
- is your app ready for PaaS?
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- installing PaaS might be easy
- operating PaaS will be tough**
- so is operating your apps tough enough to move?
- private PaaS or not?
- is your app ready for PaaS?
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- installing PaaS might be easy
- operating PaaS will be tough
- **so is operating your apps tough enough to move?**
- private PaaS or not?
- is your app ready for PaaS?
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- installing PaaS might be easy
- operating PaaS will be tough
- so is operating your apps tough enough to move?
- **private PaaS or not?**
- is your app ready for PaaS?
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- installing PaaS might be easy
- operating PaaS will be tough
- so is operating your apps tough enough to move?
- private PaaS or not?
- **is your app ready for PaaS?**
- Rainbow and Unicorn Piss:

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

- you think PaaS will solve your problems?
- it will rather clone them :)
- installing PaaS might be easy
- operating PaaS will be tough
- so is operating your apps tough enough to move?
- private PaaS or not?
- is your app ready for PaaS?
- **Rainbow and Unicorn Piss:**

<http://blog.lusis.org/blog/2014/06/14/paas-for-realists/>

To Paas or not to PaaS?

→ flynn.io

- Open source PaaS (Go)
- Uses Docker to manage containers
- Ops should be a product team, not consultants
- under development
- Git push deployment
- <https://flynn.io>

→ dokku

- The smallest PaaS implementation you've ever seen
- Docker powered mini-Heroku
- Less than 100 lines of Bash
- Git push deployment
- <https://github.com/progrium/dokku>

→ deis.io

- Open source PaaS (Python)
- Git push deployment
- On top of CoreOS w/ Docker
- <http://deis.io/>

Don't forget maestro-ng!

<https://github.com/signalfuse/maestro-ng>

To Paas or not to PaaS?

→ flynn.io

- Open source PaaS (Go)
- Uses Docker to manage containers
- Ops should be a product team, not consultants
- under development
- Git push deployment
- <https://flynn.io>

→ dokku

- The smallest PaaS implementation you've ever seen
- Docker powered mini-Heroku
- Less than 100 lines of Bash
- Git push deployment
- <https://github.com/progrium/dokku>

→ deis.io

- Open source PaaS (Python)
- Git push deployment
- On top of CoreOS w/ Docker
- <http://deis.io/>

Don't forget maestro-ng!

<https://github.com/signalfuse/maestro-ng>

To Paas or not to PaaS?

→ flynn.io

- Open source PaaS (Go)
- Uses Docker to manage containers
- Ops should be a product team, not consultants
- under development
- Git push deployment
- <https://flynn.io>

→ dokku

- The smallest PaaS implementation you've ever seen
- Docker powered mini-Heroku
- Less than 100 lines of Bash
- Git push deployment
- <https://github.com/progrium/dokku>

→ deis.io

- Open source PaaS (Python)
- Git push deployment
- On top of CoreOS w/ Docker
- <http://deis.io/>

Don't forget maestro-ng!

<https://github.com/signalfuse/maestro-ng>

To Paas or not to PaaS?

→ flynn.io

- Open source PaaS (Go)
- Uses Docker to manage containers
- Ops should be a product team, not consultants
- under development
- Git push deployment
- <https://flynn.io>

→ dokku

- The smallest PaaS implementation you've ever seen
- Docker powered mini-Heroku
- Less than 100 lines of Bash
- Git push deployment
- <https://github.com/progrium/dokku>

→ deis.io

- Open source PaaS (Python)
- Git push deployment
- On top of CoreOS w/ Docker
- <http://deis.io/>

Don't forget maestro-ng!

<https://github.com/signalfuse/maestro-ng>

To Paas or not to PaaS?

Remember, that PaaS might fail; plan & test for disaster !



Docker & CLI

```
$ docker run -t -i fedora /bin/bash  
> yum -y update  
> yum -y install nginx  
$ docker commit 73fa45674fd fedora-nginx  
$ docker run -d -p 80:80 fedora-nginx
```

Docker & CLI

Or via Dockerfile:

```
FROM fedora
```

```
MAINTAINER scollier <scollier@redhat.com>
```

```
RUN yum -y update
```

```
RUN yum -y install nginx
```

```
EXPOSE 80
```

```
CMD [ "/usr/sbin/nginx" ]
```

```
$ docker build -t fedora --rm .
```

```
$ docker run --name=nginx fedora
```

FIG?

- <http://www.fig.sh?>



Actually: Docker-Compose

- <http://docs.docker.com/compose/>
- for single host env



Docker-Compose

- <http://docs.docker.com/compose/>
- for single host env

```
FROM python:2.7
ENV PYTHONUNBUFFERED 1
RUN mkdir /code
WORKDIR /code
ADD requirements.txt /code/
RUN pip install -r requirements.txt
ADD . /code/
```

Docker-Compose

```
db:
  image: postgres
web:
  build: .
  command: python manage.py runserver 0.0.0.0:8000
  volumes:
    - /srv/app/code:/code
  ports:
    - "8000:8000"
  links:
    - db
```

Docker-Compose

```
$ alias fig='docker-compose'
```

```
$ fig run web django-admin.py startproject figexample .
```

```
$ fig up
```

Management during runtime?

```
$ fig run web python manage.py syncdb
```


Docker & Ansible

Ansible + Docker

&

Docker + Ansible

Docker & Ansible

Ansible docker core module:

http://docs.ansible.com/docker_module.html

```
- hosts: web
  sudo: yes
  tasks:
    - name: run tomcat servers
      docker: >
        image=centos
        command="service tomcat6 start"
        ports=8080
        count=5
        memory_limit=128MB
        link=mysql
        expose=8080
        registry=...
        volume=...
```

Docker & Ansible

Building image with Ansible:

```
FROM ansible/centos7-ansible:stable
ADD ansible /srv/example
WORKDIR /srv/example
RUN ansible-playbook web.yml -c local
EXPOSE 80
CMD ["/usr/sbin/nginx"]
```

Docker & Ansible

Building image with Ansible:

```
FROM ansible/centos7-ansible:stable
ADD ansible /srv/example
WORKDIR /srv/example
RUN ansible-playbook web.yml -c local
EXPOSE 80
CMD ["/usr/sbin/nginx"]
```

ansible/web.yml:

- name: Install webserver
- hosts: localhost
- tasks:
 - yum: pkg=nginx state=latest
 - shell: echo "ansible" > /usr/share/nginx/html/index.html

Docker & Ansible

Yet another demo?

CoreOS


- Designed for massive server deployments
- Support Docker container out of the box
- It's a Chrome OS fork
- Consists of couple of components:
 - SystemD – not just a init system ;)
 - Fleet – cluster level manager & scheduler
 - etcd – light & distributed key / value store
 - Docker – the only packaging method in CoreOS

CoreOS



CoreOS


Fleet – cluster level manager & scheduler



<https://coreos.com/using-coreos/clustering/>

CoreOS


etcd – light & distributed key / value store
(used for configuration store)



<https://coreos.com/docs/#cluster-management>

CoreOS


Docker – the only
packaging method
in CoreOS



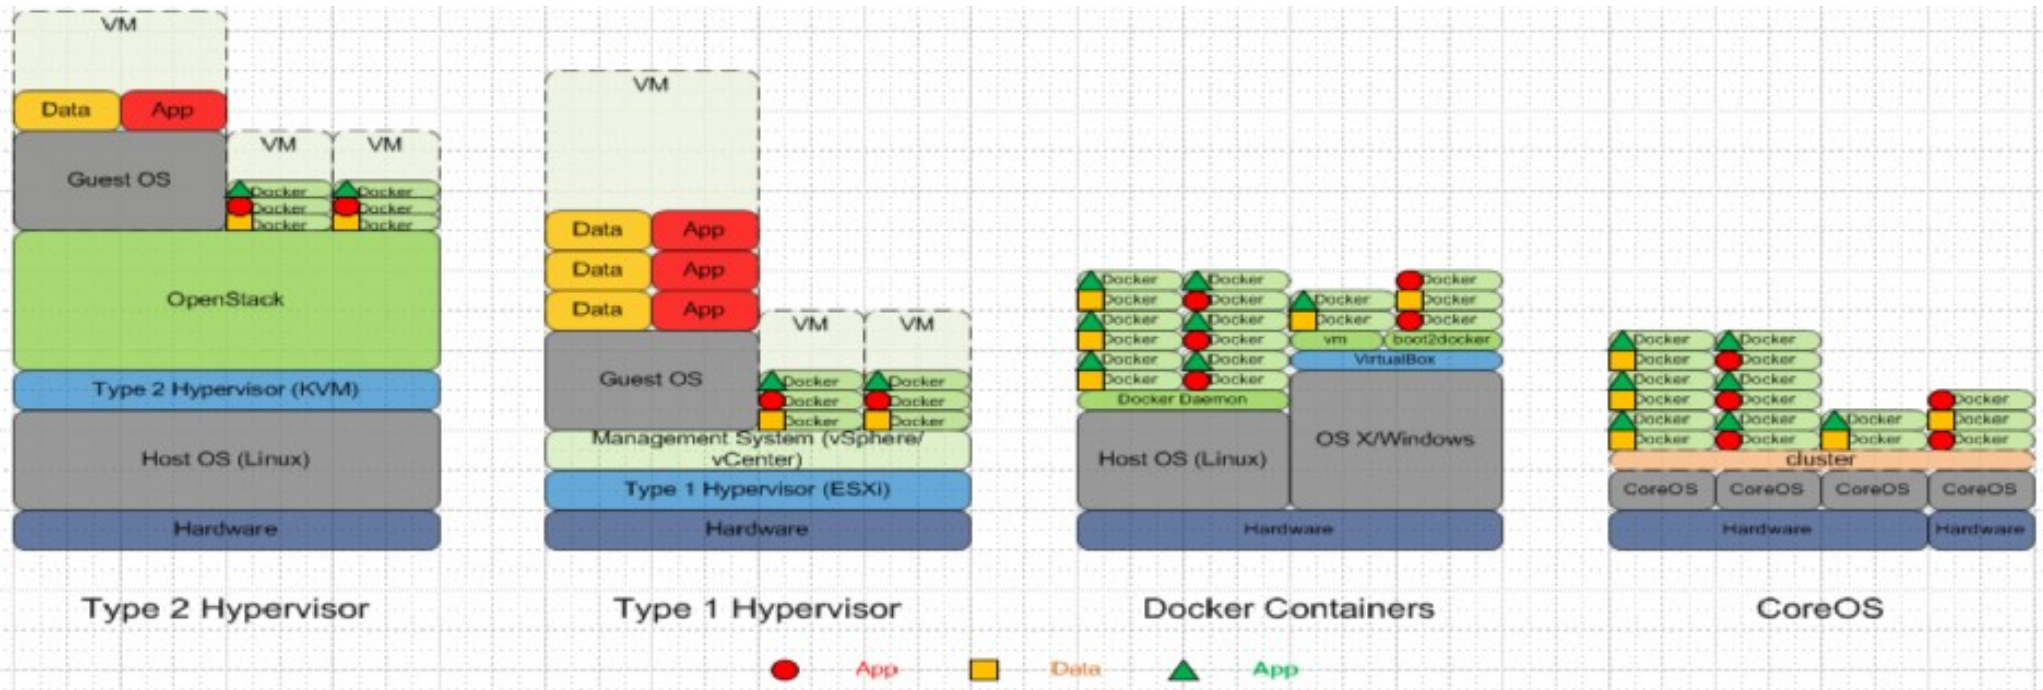
CoreOS

Full Control Over Updates:

- active-passive (dual) partitioning
- update all-packages-at-once
- safe rollback
- no more partially upgraded state
- reboot to get upgraded



CoreOS



CoreOS

Cluster management with Fleet & SystemD

[Unit]

Description=My Service

After=docker.service

[Service]

TimeoutStartSec=0

ExecStartPre=-/usr/bin/docker kill hello

ExecStartPre=-/usr/bin/docker rm hello

ExecStartPre=/usr/bin/docker pull busybox

ExecStart=/usr/bin/docker run --name hello busybox /bin/sh -c

"while true; do echo Hello World; sleep 1; done"

ExecStop=/usr/bin/docker stop hello

CoreOS

Cluster management with Fleet & SystemD

```
$ fleetctl load hello.service
```

```
Unit hello.service loaded on 8145ebb7.../10.10.1.3
```

```
$ fleetctl start hello.service
```

```
Unit hello.service launched on 8145ebb7.../10.10.1.3
```

```
$ fleetctl list-machines
```

MACHINE	IP	METADATA
148a18ff-6e95-4cd8-92da-c9de9bb90d5a	10.10.1.1	-
491586a6-508f-4583-a71d-bfc4d146e996	10.10.1.2	-
c9de9451-6a6f-1d80-b7e6-46e996bfc4d1	10.10.1.3	-

```
$ fleetctl list-units
```

UNIT	MACHINE	ACTIVE	SUB
hello.service	c9de9451.../10.10.1.3	active	running

CoreOS

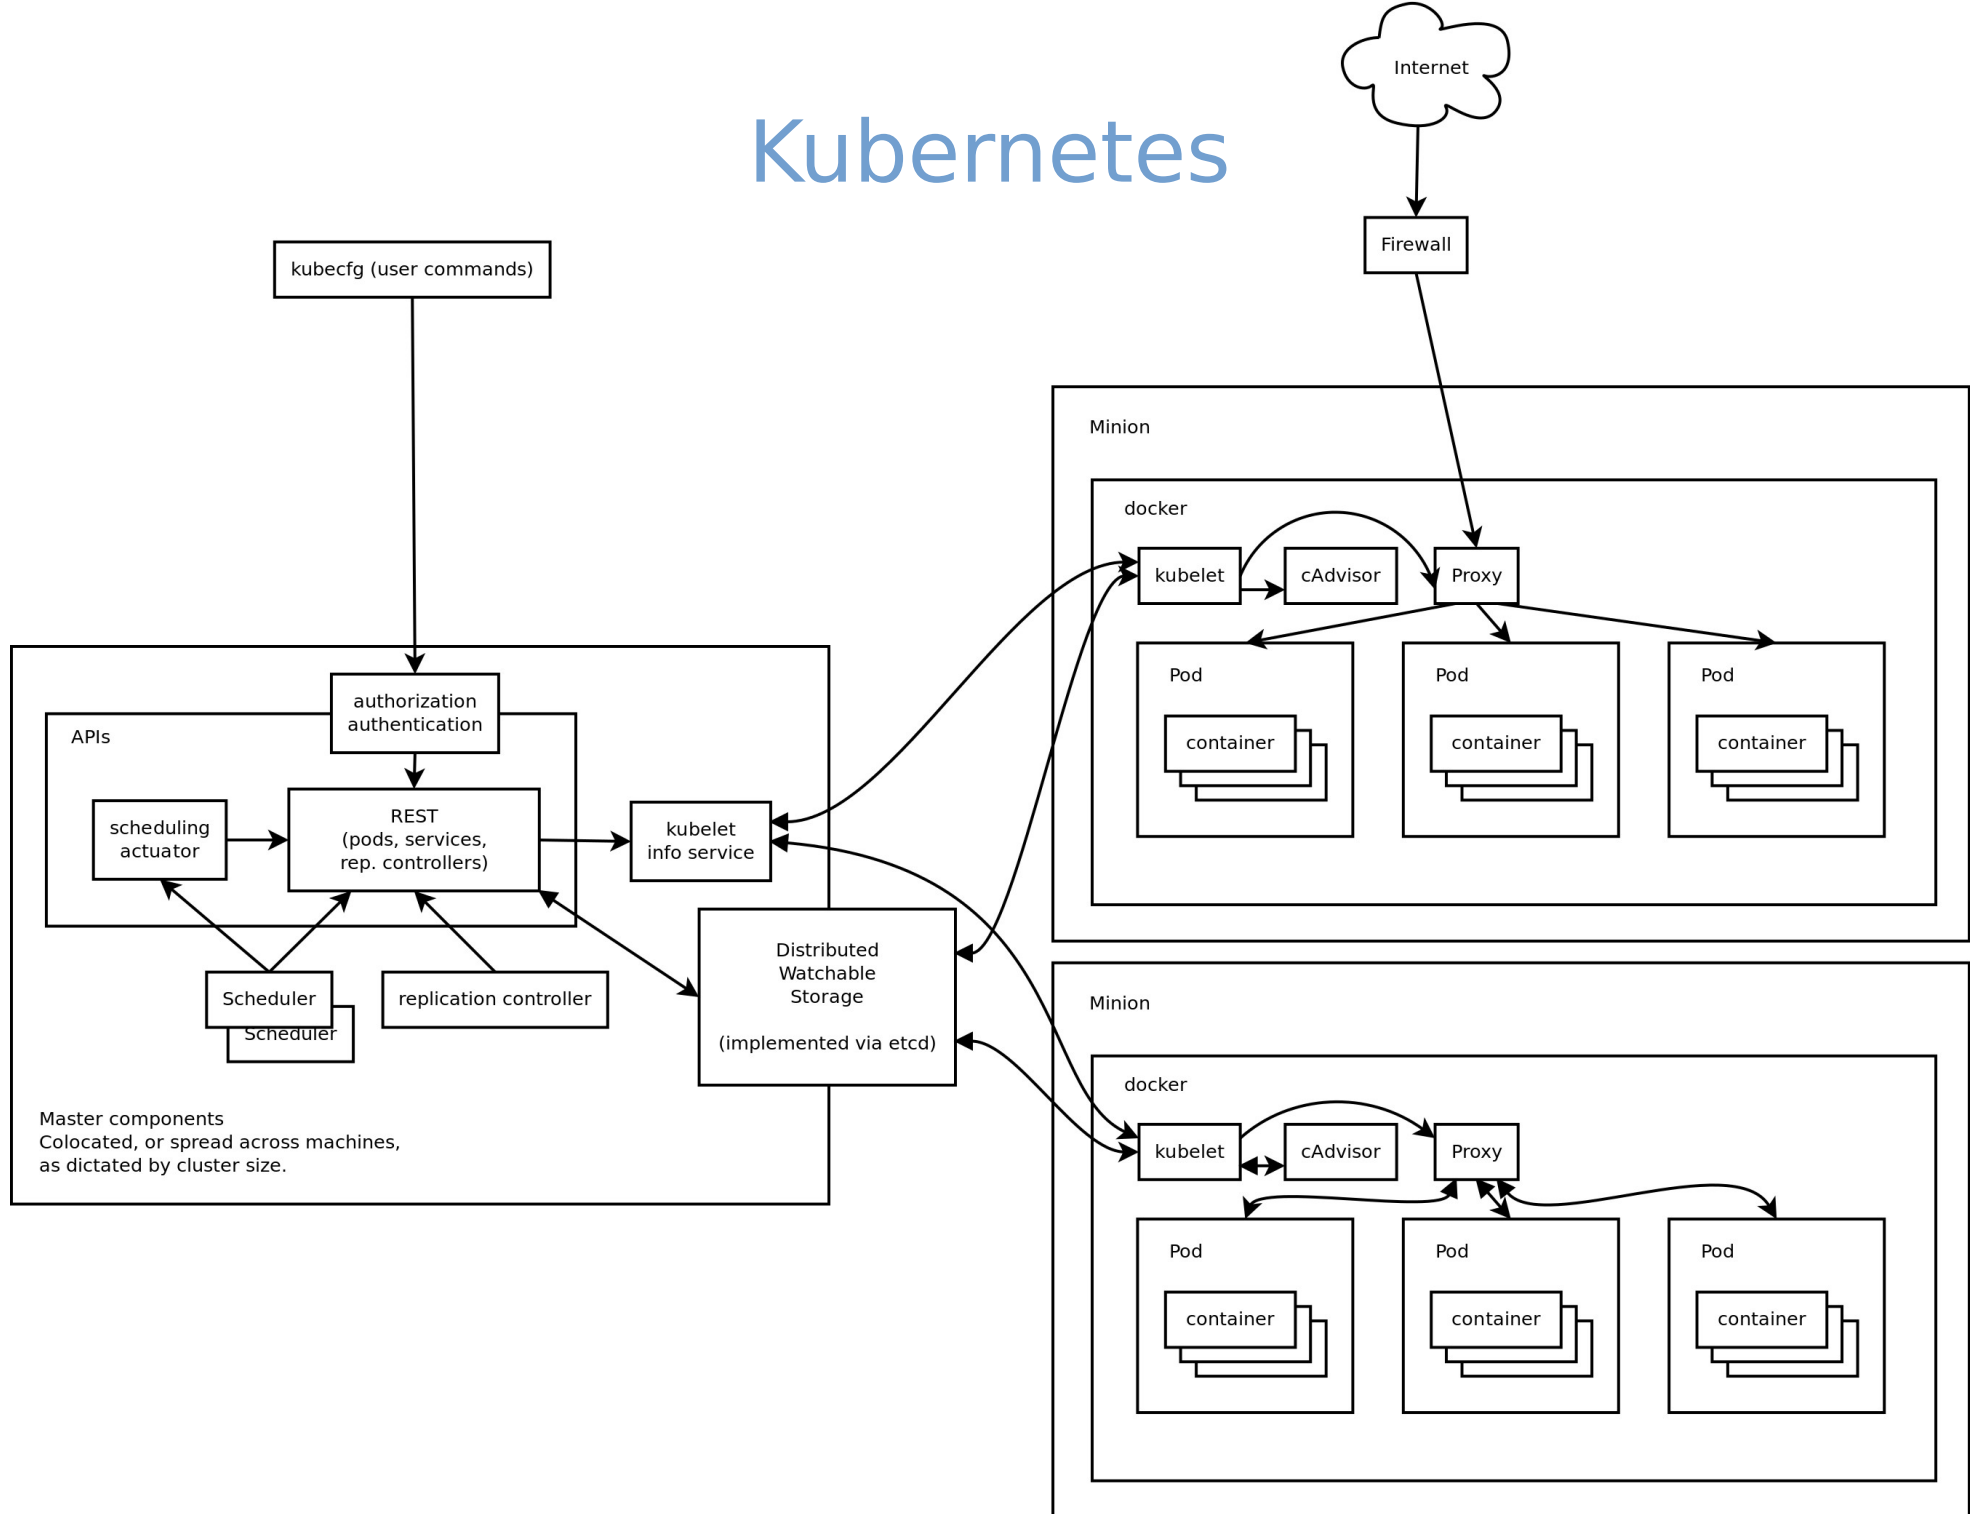
- scheduler
- HA
- dependencies

<https://coreos.com/docs/launching-containers/launching/launching-containers-fleet/>

Kubernetes

- <https://github.com/GoogleCloudPlatform/kubernetes>
- Advanced cluster manager (more than 1k hosts is a fit)
- Architecture:
 - master
 - minion
 - pod
 - replication controller
 - label

Kubernetes



Kubernetes



SmartStack

- automated service discovery and registration framework
- ideal for SOA architectures
- ideal for continuous integration & delivery
- solves “works on my machine” problem

SmartStack

- automated service discovery and registration framework
- ideal for SOA architectures
- ideal for continuous integration & delivery
- solves “works on my machine” problem


haproxy + nerve + synapse + zookeeper = smartstack

SmartStack

Synapse

- discovery service (via zookeeper or etcd)
- installed on every node
- writes haproxy configuration
- application doesn't have to be aware of this
- works same on bare / VM / docker
- <https://github.com/airbnb/nerve>

SmartStack

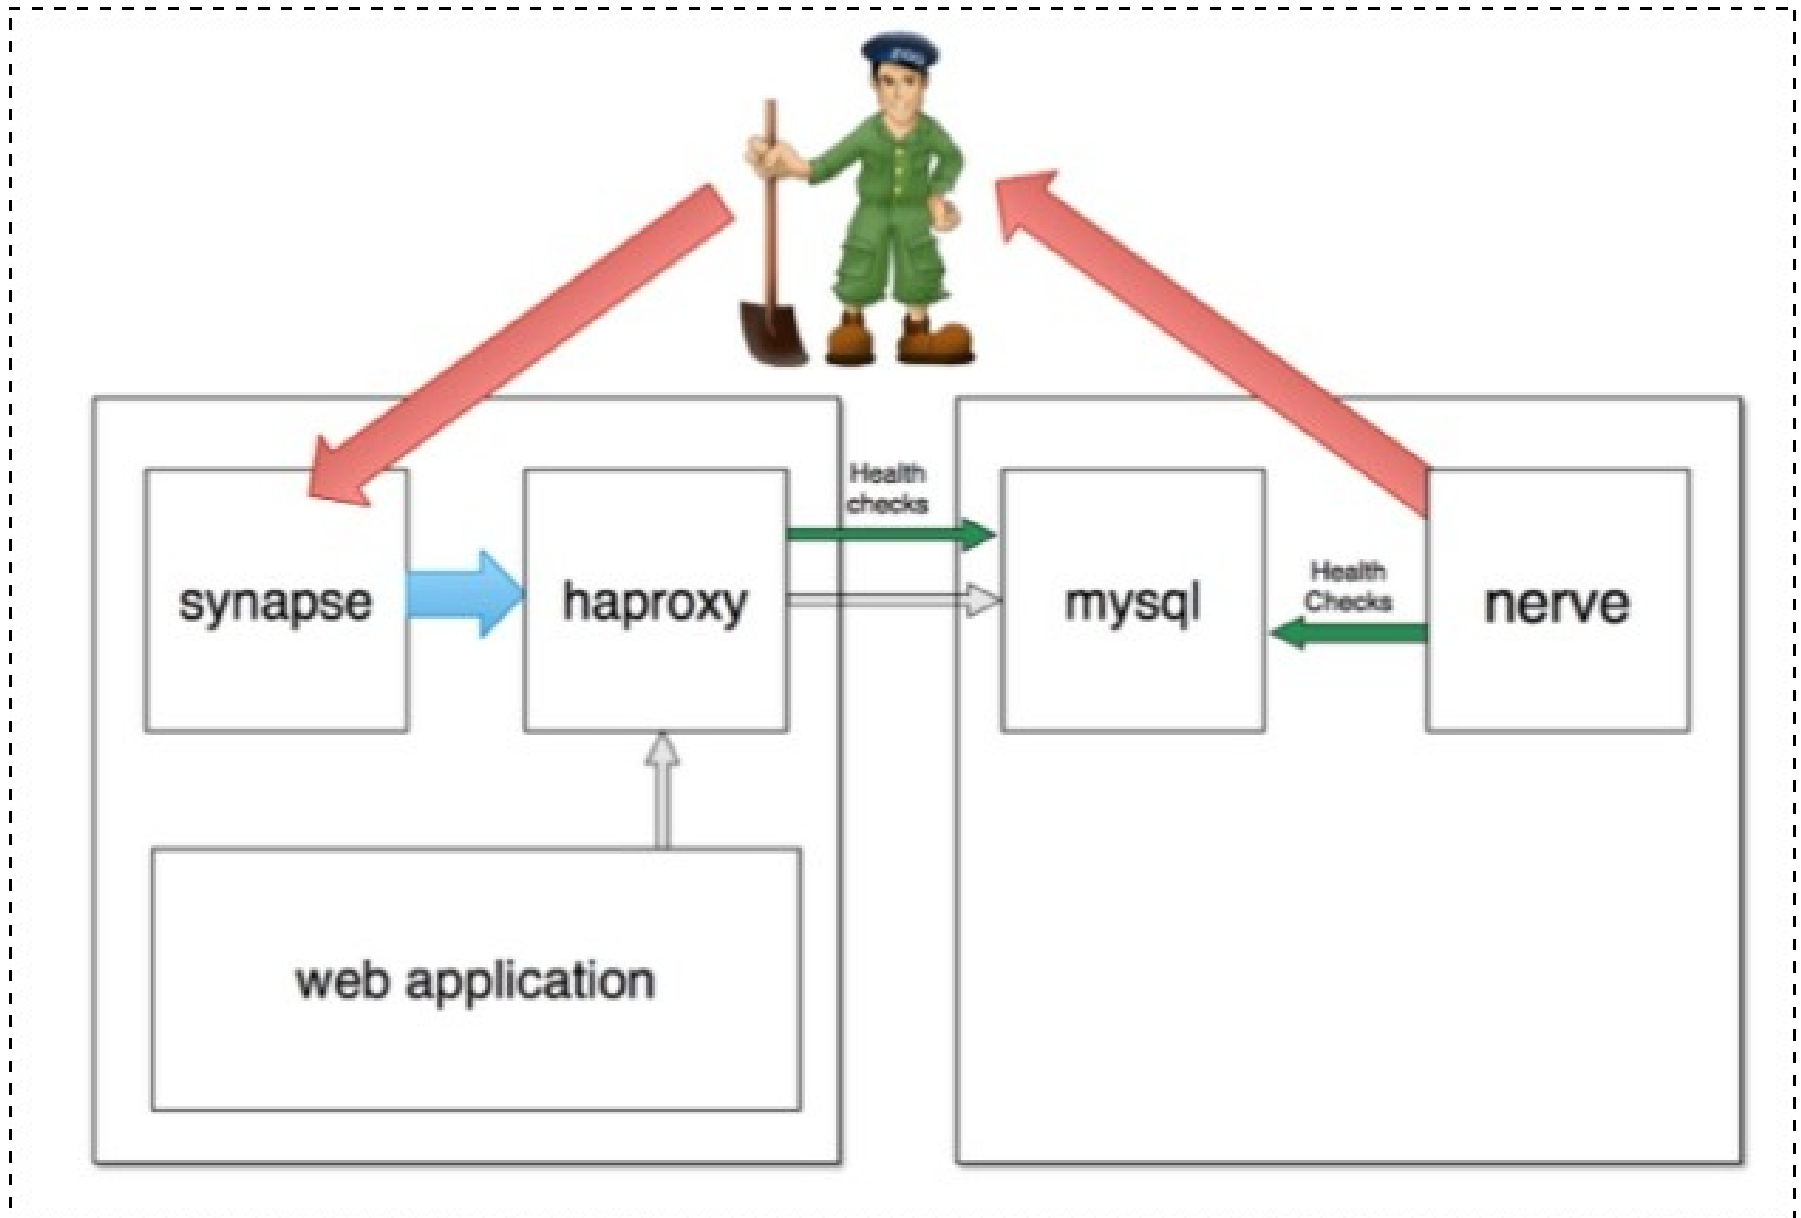


SmartStack

Nerve

- health checks (pluggable)
- register service info to zookeeper (or etcd)
- <https://github.com/airbnb/synapse>

SmartStack



SmartStack



SmartStack

Smartstack + Docker = <3

SmartStack

Smartstack + Docker = <3

but also remember about Consul
(come to #dockerkrk 2 meetup!)

Summary

Wanna learn Docker?

<http://dockerbook.com/>

Summary



 WOULD YOU LIKE TO KNOW MORE?

Freenode #docker

#KrkDocker meetups (<http://www.meetup.com/Docker-Krakow-Poland/>)

<https://github.com/docker/docker>

sources?

- [docker.io](https://docs.docker.io/) documentation
- dockerbook.com
- slideshare!
- zounds of blogposts (urls provided)
- and some experience ;)

Thank you :)

Orchestrating Docker containers at scale

Maciej Lasyk

12 Sesja Linuksowa

2015-04-18, Wrocław

<http://maciej.lasyk.info>

maciej@lasyk.info

@docent-net