

High Availability Explained

Maciej Lasyk


Kraków, devOPS meetup #2

2014-01-28

“Anything that can go wrong, will go wrong”


Murphy's law

“Anything that can go wrong, will go wrong”



Murphy's law


“Anything that can go wrong, will go wrong”



Murphy's law

An electrical explosion and fire Saturday at a Houston data center operated by The Planet has taken the entire facility offline. The company claimed power to the facility was interrupted when a transformer exploded. Official reports that three walls were blown down causing a fire.

“Anything that can go wrong, will go wrong”



Murphy's law

An electrical explosion and fire Saturday at a Houston data center operated by The Planet has taken the entire facility offline. The company claimed power to the facility was interrupted when a transformer exploded. Official reports that three walls were blown down causing a fire.

Three walls of the electrical equipment room on the first floor blew several feet from their original position, and the underground cabling that powers the first floor of H1 was destroyed.

High Availability is in the eye of the beholder

High Availability is in the eye of the beholder

CEO: we don't loose sales

High Availability is in the eye of the beholder

CEO: we don't loose sales

Sales: we can extend our offer basing on HA level

High Availability is in the eye of the beholder

CEO: we don't loose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

High Availability is in the eye of the beholder

CEO: we don't lose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

Developers: we can be proud – our services are working ;)

High Availability is in the eye of the beholder

CEO: we don't lose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

Developers: we can be proud – our services are working ;)

System engineers: we can sleep well (and fck, we love to!)

High Availability is in the eye of the beholder

CEO: we don't lose sales

Sales: we can extend our offer basing on HA level

Accounts managers: we don't upset our customers (that often)

Developers: we can be proud – our services are working ;)


System engineers: we can sleep well (and fsck, we love to!)

Technical support: no calls? Back to WoW then.. ;)

So how many 9's?

So how many 9's?

SIMPLY EXPLAINED



So how many 9's?

Monthly: 1 hour of outage means $100\% - 0.13888 \approx \mathbf{99.86112}$ of availability

So how many 9's?

Monthly: 1 hour of outage means $100\% - 0.13888 \approx \mathbf{99.86112}$ of availability

Yearly: 1 hour of outage means $100\% - 0.01142 \approx \mathbf{99.98858}$ of availability

So how many 9's?

Monthly: 1 hour of outage means $100\% - 0.13888 \approx \mathbf{99.86112}$ of availability

Yearly: 1 hour of outage means $100\% - 0.01142 \approx \mathbf{99.98858}$ of availability

Availability	Downtime (year)	Downtime (month)
90% ("one nine")	36.5 days	72 hours
95%	18.25 days	36 hours
97%	10.96 days	21.6 hours
98%	7.30 days	14.4 hours
99% ("two nines")	3.65 days	7.2 hours
99.5%	1.83 days	3.6 hours
99.8%	17.52 hours	86.23 minutes
99.9% ("three nines")	4.38 hours	21.56 minutes
99.99 ("four nines")	52.56 minutes	4.32 minutes
99.999 ("five nines")	5.26 minutes	25.9 seconds

So how many 9's?




Figure 1-3 Cost of availability as opposed to loss due to availability

<https://jazz.net/wiki/bin/view/Deployment/HighAvailability>

HA terminology

RPO: Recovery Point Objective; how much data can we loose?

HA terminology

RPO: Recovery Point Objective; how much data can we lose?

RTO: Recovery Time Objective; how long does it take to recover?


HA terminology

RPO: Recovery Point Objective; how much data can we loose?

RTO: Recovery Time Objective; how long does it take to recover?

MTBF: Mean-Times-Between-Failures; time between failures

(density fnc -> reliability fnc) $\int_0^{\infty} f(t) dt = 1.$



Time Between Failures = { down time - up time }

https://en.wikipedia.org/wiki/Mean_time_between_failures

HA terminology

SLA: Service Level Agreement;
formal definitions (customer <-> provider)

HA terminology

SLA: Service Level Agreement;
formal definitions (customer <-> provider)

OLA: Operational Level Agreement; definitions within organization;
help us keeping provided SLAs

SLAs..

So what is written in **SLAs**?

Availability	Downtime (year)	Downtime (month)
90%	36.5 days	72 hours
95%	18.25 days	36 hours
97%	10.96 days	21.6 hours
98%	7.30 days	14.4 hours
99%	3.65 days	7.2 hours
99.5% (EC2, EBS)	1.83 days	3.6 hours
99.8%	17.52 hours	86.23 minutes
99.9% (SoftLayer, IBM)	4.38 hours	21.56 minutes
99.99	52.56 minutes	4.32 minutes
99.999	5.26 minutes	25.9 seconds

SLAs..

So what is written in **SLAs**?

Availability	Downtime (year)	Downtime (month)
90%	36.5 days	72 hours
95%	18.25 days	36 hours
97%	10.96 days	21.6 hours
98%	7.30 days	14.4 hours
99%	3.65 days	7.2 hours
99.5% (EC2, EBS)	1.83 days	3.6 hours
99.8%	17.52 hours	86.23 minutes
99.9% (SoftLayer, IBM)	4.38 hours	21.56 minutes
99.99	52.56 minutes	4.32 minutes
99.999	5.26 minutes	25.9 seconds

<http://aws.amazon.com/ec2/sla/>

<http://www.softlayer.com/about/service-level-agreement>

SLAs..

Availability mentioned in **SLAs** are only **goals** of service provider

Usually when it's **not met** than company **pays off** the fees

How deep is this hole?

app layer (core, db, cache)

data storage

operating system

hardware

networking

location

So we would like to achieve 99,9999% which is about 30s of downtime per year

How deep is this hole?

app layer (core, db, cache)

data storage

operating system

hardware


networking

location

Even Proof of Concept is very hard to provide: 5s of downtime per layer yearly!

Load-balancing and failover


LB:



<http://www.netdigix.com/linux-loadbalancing.php>

Load-balancing and failover

Failover:



<http://www.simplefailover.com/>

LB – 4th layer or 7th?

4th layer:

- high performance
- just do the LB work!
- reliable
- scalable


7th layer:

- low cost
- good for quickfixes / patches
- not that scalable
- low performance
- complex codebase
- custom code for protocols
- cookies? what about memcache..

Disaster Recovery




Disaster Recovery



<http://disasterrecovery.starwindsoftware.com/planning-disaster-recovery-for-virtualized-environments>

Disaster Recovery



<http://disasterrecovery.starwindsoftware.com/planning-disaster-recovery-for-virtualized-environments>


Hot site: active synchronization, could be serving services. Cost can be high

Warm site: periodical synchronization, DR tests needed. Low costs

Cold site: Nothing here – just echo and some place to spin services; nightmare

Planning for failure

ENTEPRISE ARCHITECTURE MADE EASY



Planning for failure

Everything starts here - DNS:

- keep TTLs low (300s). Can't make under 60min? That's bad!
- check SLA of DNS servers (dnsmadeeasy.com history)
- what do you know about DNSes?
 - zero downtime here is a must!
 - this can be achieved with complicated network abracadabra
 - remember what 99.9999% means?
- round robin is a load – balancer but **without failover!**
- GSLB – killed by OS/browser/srvs cache'ing
(GlobalServerLoadBalancing)
- GlobalIP (SoftLayer etc) – workaround for GSLB via routing

Planning for failure

E-mail servers:

- it's simple as MX records (**delivering**)
- it's almost simple as complicated system of SMTP servers (**sending**)
- it's not that simple when IMAP locking over DFS (**reading**)

```
5 gmail-smtp-in.l.google.com.  
10 alt1.gmail-smtp-in.l.google.com.  
20 alt2.gmail-smtp-in.l.google.com.  
30 alt3.gmail-smtp-in.l.google.com.  
40 alt4.gmail-smtp-in.l.google.com.
```

When MXing – **watch the spam!**

Planning for failure

WEB servers:

- it's simple as some frontend loadbalancer
- did you **really** stick user session to particular server? Memcache!
- LB balancing algorithm
- how many Lbs?
- what if LB goes down?

Planning for failure

DB servers:

- it's.. not that simple
- replication (master – master? App should be aware..)
- replication ring? Complicated, works, but in case of failure...
- let's talk about MySQL:
 - NoSPOF solution: MySQL cluster
 - MySQL Galera cluster – synch, active-active multi-master
 - master – master – simply works
 - Failover? Matsunobu Yoshinori `mysql-master-ha`
 - MySQL utilities (<http://www.clusterdb.com/mysql/mysql-utilities-webinar-qa-replay-now-available/>)

Planning for failure

Caching servers:

- this is cache for God's sake – why would we use HA here?
- just use proper architecture like... redundancy.

Planning for failure

Caching servers:

- this is cache for God's sake – why would we use HA here?
- just use proper architecture like... redundancy.

Load – balancers:

- remember about failovering IP addresses!

Planning for failure

Caching servers:

- this is cache for God's sake – why would we use HA here?
- just use proper architecture like... redundancy.

Load – balancers:

- remember about failovering IP addresses!

Storage – DFSes:

- GlusterFS – we'll see it in action in a minute
- NFS? Could be – over some SAN / NAS (high cost solution)
- CephFS – just like GlusterFS – it's great and does the work
- DRBD – lower level, does the work on block – device layer – slow...

Planning for failure

GlusterFS:

- low cost (could be..)
- distributed volumes
- replicated volumes
- striped volumes
- and...
 - distributed – striped volumes
 - distributed – replicated volumes
 - distributed – striped – replicated volumes
- sound good? :)

Planning for failure

GlusterFS: replicated volumes vs Geo-replication

- replicated:

- mirrors data
- provides HA
- synch – replication

- Geo-replication:

- mirrors data across geo – distributed clusters
- ensures backing up data for DR
- asynch – replica (periodic checks)

Planning for failure


HA for virtualization solutions?

- it's really complicated, like...

Planning for failure

HA for virtualization solutions?

- it's really complicated, like...





Tools

The most important tool would be the conclusion from the picture below:

Tools



The most important tool would be the conclusion from the picture below:



LET THE CLOUDS MAKE YOUR LIFE EASIER

Tools

The most important tool would be the conclusion from the picture below:



LET THE CLOUDS MAKE YOUR LIFE EASIER



Tools

- DNS: roundrobin, GSLB, low ttls, globalIP

Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx

Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx
- Failover (statefull services):
 - IP: KeepAlived + sysctl

Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx
- Failover (statefull services):
 - IP: KeepAlived + sysctl
 - Managing: pacemaker (manager) + corosync (message'ing)

Tools

- DNS: roundrobin, GSLB, low ttls, globalIP
- Load-Balancers (I7, stateless services): HaProxy, Pound, Nginx
- Failover (statefull services):
 - IP: KeepAlived + sysctl
 - Managing: pacemaker (manager) + corosync (message'ing)
- (almost) All-In-One: Linux Virtual Server

Turn on HA thinking!

Main goal of HA? **Improve user experience!**

- keep the app **fully functional**
- keep the app **resistant and tolerant** to faults
- provide method for a **successful audit**
- **sleep well** (anyone awake?) ;)

Thank you :)

High Availability Explained

Maciej Lasyk

Kraków, devOPS meetup #2

2014-01-28

<http://maciek.lasyk.info/sysop>

maciek@lasyk.info

@docent-net